

Neural Architecture Search and tinyML: A Survey

Daniel Duclos-Cavalcanti

Abstract—There is no denying the increasing success that Deep Neural Networks (DNNs) have displayed across various tasks such as image classification, speech recognition, machine translation and many others. This progress can be attributed largely due to *architecture engineering*, which is a process that requires immense domain expertise, intuition and some amount of trial and error. This is not ideal as it is both time consuming and error-prone. Neural Architecture Search (NAS) rises as the next logical step aiming to both automate architecture discovery and further the understanding of the inner workings of DNNs. This field has grown remarkably within the last 5 years and among the many challenges of NAS, there are concerns regarding computational costs and time feasibility. This however changes within the context of *tinyML*, which is an expanding field at the intersection of machine learning and embedded systems. Due to the resource constrained conditions involved in most embedded devices, various techniques have been developed to reduce size and memory consumption of deep learning models by the tinyML community. Due to that, the stages concerning training and performance evaluation become substantially faster in comparison to the models used in NAS research. Thus, presenting an interesting opportunity to explore NAS applications within a different paradigm. Throughout this work, an overview of existing research in NAS, specifically concerned with the use of evolutionary algorithms methods will be presented, as well as highlighting relevant applicabilities to tinyML.

Index Terms—Neural Architecture Search (NAS), evolutionary computation (EC), Deep Learning, tinyML.

I. INTRODUCTION

The advancement of deep learning, although extremely beneficial, has also caused a continuous demand for architecture design. This coupled with a growing model complexity, demands ample time and expert knowledge to either improve existing architectures or create new ones for novel applications or problems.

After the work in [1] proposed by Google, it was shown for the first time that NAS algorithms have the potential to find models that rival the current state of the art. However, done so in an automated fashion, minimizing human participation. Since then, many different methods have appeared.

To better understand NAS, one can categorize it using the following characteristics [2]:

- **Search Space:** What is the set of all possible architectures that can in principle be considered by the algorithm.
- **Search Strategy:** How the *search space* is explored by the algorithm.
- **Performance Estimation Strategy:** How performance is evaluated at every architecture iteration.

A. Search Space

NAS is an optimization problem, whose search space is the defining factor to its complexity. The smaller the search space is, the faster the search may converge, as well as requiring less

computational resources. However, this limits the freedom to explore architectures that have never been conceived before, and also possibly restricts the complexity of the design.

How an architecture is constructed or described is a core concept to how the search space of an NAS algorithm is defined. The simplest approach would be to define an architecture as a *chain-structured neural network*, which essentially consists of a sequence of layers whose inputs are the output of their preceding layer. In this case the space is parametrized by maximum number of layers, type of operations per layer, and hyperparameters conditioned by the chosen operation. An illustration to this can be seen to the left of Fig. 1.

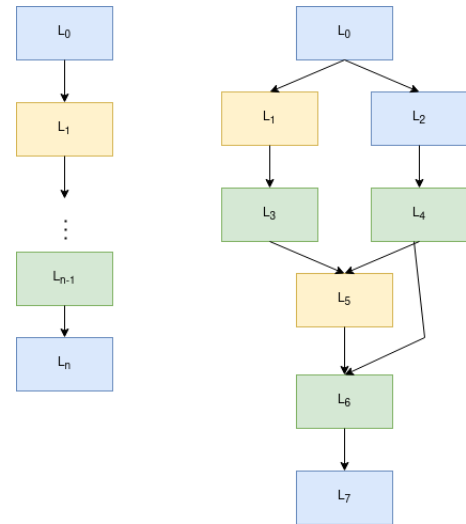


Fig. 1: Illustration of different ways to model neural architectures. Left: an example of a chain-structured NN. Right: an example of a multi-branch NN.

The next step would consist of including multiple layer connections and other modern design elements such as skip connections, which has already been seen in [1] and [3]. This allows to build more complex *multi-branch networks*, which cannot be described as simple sequential layer chaining, but as a structure where each layer's input is a function of previous layers outputs. This increases significantly the degrees of freedom of architecture design, which leads to a much larger search space. An illustration to this can once again be seen to the right of Fig. 1. A chain-structured neural network can therefore be defined as a special case of the multi-branch network.

A trend is however the search within *cells*, initially considered by works such as [4] and [5]. What is proposed, is to break-off architectures into cells, such that the search space is then designated within a single cell instead of an entire architecture. This cell is the basic unit of the resulting

architecture, which can be then combined with other identical cells in various fashions to produce a final design. This strongly reduces the search space, as there are substantially less layers within a single cell in comparison to an entire architecture.

Cells drive the discussion between *micro-architectures* versus *macro-architectures* [2]. The macro architecture attempts to determine how cells should be connected and how many are needed to build a model [2]. On the other hand, the micro architecture aims to find the optimal structure within a cell [2]. Ideally, both should be optimized simultaneously, which of course leads to a complex search space. There have been efforts to minimize this endeavor by fixing macro-architectures with known working topographies such as in [6] with DenseNet [7]. This practice, dubbed as *human knowledge injection* attempts to reduce the search space through applying domain expertise known to obtain good results.

Another trend, predominantly used, is the utilization of *blocks* as the basic units of the search space. Here, various layers of different types are combined to produce basic blocks such as ResBlock, DenseBlock, ConvBlock and so on. These blocks are fixed in their topology and require less parameters to build an architecture. Therefore, they drastically reduce the search space. Various blocks can be combined in different manners and quantities to produce competitive architectures. Cells can be considered as a special case of block-based search, where each block is identical [8].

The same modelling used on multi-branch networks can be used with cells or blocks, simply replacing layers with these units.

B. Search Strategy

As any space search problem, there is a *exploration-exploitation* trade-off to be considered [2]. Obtaining a high-performing architecture within a feasible amount of time is desired. However, converging too early to a suboptimal result is also not the goal.

Based on the current state of the art, NAS search algorithms can be classified mainly into three different categories [8]:

1) **Reinforcement Learning (RL)-based NAS Algorithms:**

Here, one considers the development of a neural architecture as the agent's action, where the action space is the same as the search space. Therefore, it is then possible to frame NAS as a reinforcement learning problem [2]. After evaluating the performance of the given trained architecture on fitness data, it is possible to determine the agent's reward. How this estimation is performed will depend on the given method, more to these approaches will be seen at section I-C. Furthermore, how the agent's policy is represented and its optimization will also vary. More on trade-offs and details to these approaches will be seen in II.

2) **Evolutionary Computation (EC)-based NAS Algorithms:**

By the application of well established EC methods, which are various techniques based on the evolution of species within nature, many different efforts were implemented to navigate their respective search space.

3) **Gradient-based NAS Algorithms:** EC and Reinforcement learning methods consider an approach, where neural architectures are spread across a discrete and non-differentiable search space. There are many reasons for this, as networks may vary through many different variables other than numerical parameters or weights. Gradient-based approaches however impose restrictions and reformulate the problem to obtain a continuous search space. This allows an algorithm to use gradient-based methods to pursue an optimal model, which leads many times to competitive and fast techniques.

More on these different techniques will be seen in section II-B.

C. Performance Estimation Strategy

Independent of the search strategy, it is necessary to know how any given architecture performed in order to guide the next steps of the algorithm. There are many ways to estimate said performance, whereas the simplest would be complete training and validation. Given the complexity and size of the search spaces within NAS, this requires GPU days in the order of thousands [2]. This is why extensive research has been employed to reduce time on performance estimation, since it is a significant time bottleneck. Some of these approaches include [2]:

- **Lower Fidelity Techniques:** Shorter training time [9], [5], training solely on a subset of the data [10], training on downsampled data [11] or with downsampled models [5], [3]. These methods do introduce bias as performances will normally be underestimated.
- **Learning Curve Extrapolation:** Performance is extrapolated after just a small number of epochs and then decided upon directly. Klein et al. [10] considered architectural hyperparameters to predict which architectures are most promising after partial learning. Domhan et al [12] extrapolated partial learning curves to predict and eliminate sub-optimal architectures. The difficulty here lies in generating good predictions within a small amount of evaluations, in order to speed up the search.
- **Weight Inheritance:** also dubbed as *network morphisms* is a technique that passes down weights from previously trained models to new ones. This approach can cut down computational costs to just a few GPU days [13].
- **One-Shot Models:** also called *weight sharing* is a technique that treats all architectures as subgraphs of a supergraph, named the one-shot model. Weights are shared between architectures that meet the condition of having edges in common. Finally, only the weights of a single one-shot model has to be trained and the sub-graph architectures can be evaluated directly as they inherit weights from the one-shot model. This demonstrates great success, as it was one of the main reasons that the DARTS algorithm was able to cut down computational costs to a few GPU days [14]. However, significant bias is introduced as the underestimation of architectures by this approach is harsh.

II. STATE OF THE ART

In the following sections, the state-of-the-art of NAS and tinyML-NAS crossovers will be discussed. In section II-A, a general overview of the main NAS approaches will be presented, as well as their contrasting qualities. Then, in section II-B a deeper dive in *Evolutionary Neural Architecture Search* (ENAS) algorithms is presented. Finally, in section II-C, relevant work that brings NAS to tinyML applications will be illustrated.

A. Overview

Currently RL-based algorithms are extremely costly in terms of computation, requiring hundreds of GPUs for days. This is already the case for median-scale data sets, so as data grows more complex, so does the processing time needed to search for a suitable architecture.

Gradient-based algorithms are faster, there are examples such as the DARTS algorithm, where processing power is cut down to single digit GPUs [14]. However, due to the not completely compatible relationship of NAS and gradient-based optimization, frequently sub-optimal architectures are found [8].

On the other hand, EC methods, while not perfect, have been around for decades and are easily applicable to solve complex non-convex optimization problems [15], as they are insensitive to local minima and do not require gradient information [8].

B. Evolutionary Neural Architecture Search (ENAS)

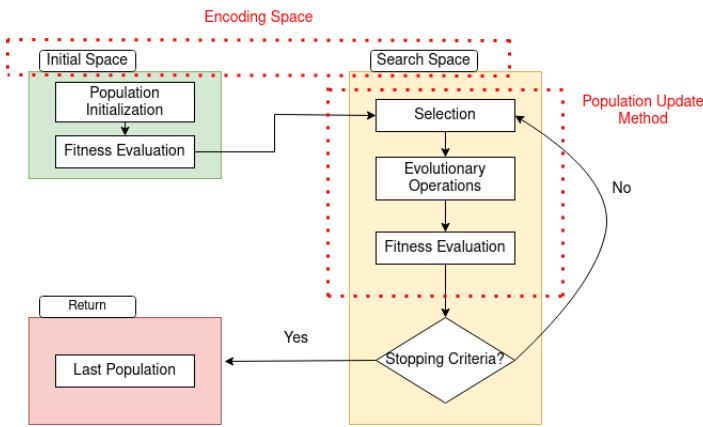


Fig. 2: Generic flowchart of a typical ENAS algorithm

ENAS algorithms are NAS algorithms that leverage existing evolutionary computational methods to search for an optimal architecture within a defined search space. It can be broken down into a series of steps exemplified by Fig. 2. The term *population* seen in Fig. 2 is defined as a finite number of *individuals*, where each individual is an iteration of a single neural architecture within the defined search space.

At the beginning of the algorithm an initial population is created within a pre-defined initial space. Then, their fitnesses are evaluated, which essentially is the process of performance estimation already covered in I-C for each individual in this generation's population. Finally, this population undergoes a

repetitive process that includes *selection, evolutionary operations* and once again a *fitness evaluation*. If the resulting *population* meet the defined goal, the final architecture has been found. Otherwise repeat the last process until this goal is reached.

Based on the main steps of an ENAS implementation, it is possible to define ENAS algorithms through four dimensions [8]: Encoding Space, Encoding Strategy, Population Update Method and Fitness Evaluation.

In the following subsections, state-of-the-art ENAS methods will be discussed and categorized by the different elements that compose an ENAS algorithm. Although, ENAS encompasses many different techniques such as evolutionary algorithms (EA), Swarm Intelligence (SI) methods and many others; this work will focus mostly on EA, since the vast majority of relevant ENAS research has utilized it.

1) *Encoding Space*: The encoding space contains all the valid individuals or architectures encoded in the population [8]. It can be further divided into the initial space and the search space. These two may be the same in some cases, but often are not. The initial space is the set of all possible architectures that any individual in the initial population may become. There are three types of architecture initialization approaches [8]:

- **Starting from trivial conditions**: This method aims to impose little to no human bias in the initial population. Real *et al.* [3] had the benefit of giving much more freedom to the algorithm to explore unseen architectures and also justified well the used of EC-methods instead of any other. However, this comes at the cost of high computational resources. The work from Xie and Yuille [16] is also another example of competitive architecture achievements that used a trivial initial space.
- **Rich Initialization**: This technique, also named *well-designed space*, consists of initializing the population within a set of state-of-the-art architectures. This way, a good architecture can be found early on and potentially decrease search time. On the other hand, finding novel architectures is very improbable. There has been relevant work that used this method as initialization technique, such as from [17].
- **Random Initialization**: This form of initialization also aims to reduce human bias or intervention in the process of architecture evolution. Many efforts have utilized this method such as [18] and [19]. Here, the search space and the initial space are the same.

The search space and their constraints have been touched in section I-A, where the concept of block-based and cell-based search spaces have been introduced. Notable works like [19] use traditional blocks such as ResBlock, DenseBlock and ConvBlock to further constraint the search space and facilitate architecture discovery. Other efforts such as [20] and [21] used different basic blocks to adapt to their respective desired architecture use-cases.

Most implementations of cell-based searches has the macro-architecture be determined by human expertise and concentrated mostly on optimizing their micro-architectures. This was done by [22].

Nonetheless, there are more EA specific parameters to consider when evaluating the search space of an ENAS algorithm. Evolutionary operators play a role in constraining the algorithm's search space and will be touched on in more depth in II-B3. An example would be the work of Irwin-Harris *et al.* [23], which did not specify explicitly the maximum depth of the resulting architectures and then used the algorithm's evolutionary operators to extend the individual's architectures indefinitely if needed.

2) *Encoding Strategy*: The encoding strategy defines how network architectures are encoded into individuals. They can be broadly defined in *fixed-length* and *variable-length* encoding strategies. These approaches differ in whether the length of an individual varies during the evolutionary process.

In the fixed-length variant, all individuals have the same length during the search process. This facilitates the process, since standard evolutionary operations are by design to be used individuals with equal lengths. During the work of Xie *et al.* [16], this enabled a much easier implementation of its ENAS algorithm, especially with the utilization of crossover. However, for fixed-length implementations, a maximum length has to be established, which then also creates a dependency on human expertise.

On the other hand, the variable-length method does not require human knowledge regarding layer depth, which creates a potential to have a fully-automated implementation. Moreover, the greatest advantage to this approach is that the algorithm can define more details of the architecture with freedom. This comes at the cost of a more difficult and complex implementation, since evolutionary operators may not be directly compatible with this encoding strategy, requiring a redesign. In addition to that, the extra freedom provided by this technique may also cause overly complex architectures with costly performance evaluations. The work of [18] is a prime example of a variable-length encoding strategy that needed to redesign its operators.

3) *Population Update Method*: The population update method includes selection strategies and how evolutionary operators are applied. This will then define how populations are ultimately updated.

Genetic Algorithm (GA)-based ENAS is the most popular approach among ENAS methods, since architecture representation is very convenient in GA [8]. Selection is the first stage of updating the current population and can be divided into several strategies, four of whom are widely used: Elitism, Discard worst or oldest, Roulette and **Tournament Selection**.

Works such as [24] take use of elitism, where essentially only the fittest of individuals are selected to compose the next population. This can cause a loss of diversity, as there is a chance that similar architectures perform similarly well and thus may cause generations to breed only akin individuals, since only the fittest are kept. This in turn, can cause the population to fall within a local optima and not being able to explore efficiently.

Other approaches such as [3] and [25] discarded the oldest individual from the population, which is also known as aging evolution. This ensures that the search does not focus on good models too early and therefore performs a more broad search

of the encoding space in comparison to non-aging evolution.

There is also the possibility to combine different strategies such as in [26], where both discarding the worst and the oldest were used.

C. NAS and tinyML

Machine learning on small microcontrollers is a great ambition. TinyML is the field that answers to that challenge and aims to provide intelligent features to even the cheapest off-the-shelf microcontrollers. This is not a small effort as microcontrollers possess very limited resources, especially concerning memory and storage. In the following subsections, relevant work that combined NAS algorithms within the tinyML universe will be shown.

1) *MCUNet*: MCUNet is a framework that shows promising results, composed of two major components: the efficient neural architecture (TinyNAS) and the lightweight inference engine (TinyEngine) [27]. The TinyNAS algorithm is a two-stage neural architecture search method, where firstly it optimizes the search space according to the given resource constraints and then performs an ENAS algorithm to find the best architecture within this new search space. Since the performance of NAS methods depend strongly on the search space [28], this technique through extra constraints forces the search to consider a smaller set, where only architectures that fit the desired requirements can be found. These requirements may be limited memory consumption, storage limits, latency and even energy.

MCUNet used weight sharing [27], where a single super network is created, which contains all the sub-networks in the search space. This method leverages the graph attribute of these networks to share weights among subnetworks and drastically reduce the time needed to train all possible individuals.

The ENAS algorithm used in MCUNet uses elitism as its selection strategy, always choosing the top-20 individuals in terms of accuracy within a given generation. Crossover is applied to generate 50 new candidates and then mutation with a probability of 10% is used to generate the remaining 50. All generations are of size 100. Finally, after 30 iterations, the fittest architecture is chosen.

In addition to that, TinyEngine implements a code generator-based compilation method that not only eliminates memory overhead, but also improves the speed of inference as well [27]. This addresses the shortcomings of similar libraries such as TF-Lite Micro [29] and CMSIS-NN [30] that opt for runtime interpretation of code instead.

Finally, this work has shown to be able to reduce memory usage by 2.7x and improve inference speed by 1.7-2.2 compared to TF-Lite Micro and CMSIS-NN, also decreased code size by up to 4.5X and 5.0x for TF-Lite Micro and CMSIS-NN respectively [27]. MCUNet achieved state of the art performance, taking 12.5 GPU days to design a model [27]. This is a great improvement compared to MnasNET, which took 40,000 GPU hours for the same data set [31] and is also faster than most NAS methods performed on regular machines.

III. CONCLUSION

Throughout this work a brief overview of the current research on NAS has been displayed. NAS is not only a complex optimization problem, but also one that requires significant computational resources in most applications. The trade-offs of the mainstream approaches to NAS have also been demonstrated, whilst also highlighting efforts done with the use of EC-based approaches.

In addition to that, different applications of NAS methods within the tinyML paradigm have been illustrated. For IoT devices and microcontrollers it is not only needed to train models considering performance, but also considering other parameters such as latency, memory usage and even energy preservation.

REFERENCES

- [1] B. Zoph and Q. V. Le, "Neural architecture search with reinforcement learning," *arXiv preprint arXiv:1611.01578*, 2016.
- [2] T. Elsken, J. H. Metzen, and F. Hutter, "Neural architecture search: A survey," *The Journal of Machine Learning Research*, vol. 20, no. 1, pp. 1997–2017, 2019.
- [3] E. Real, S. Moore, A. Selle, S. Saxena, Y. L. Suematsu, J. Tan, Q. V. Le, and A. Kurakin, "Large-scale evolution of image classifiers," in *Proceedings of the 34th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, D. Precup and Y. W. Teh, Eds., vol. 70. PMLR, 06–11 Aug 2017, pp. 2902–2911. [Online]. Available: <https://proceedings.mlr.press/v70/real17a.html>
- [4] Z. Zhong, J. Yan, W. Wu, J. Shao, and C.-L. Liu, "Practical block-wise neural network architecture generation," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 2423–2432.
- [5] B. Zoph, V. Vasudevan, J. Shlens, and Q. V. Le, "Learning transferable architectures for scalable image recognition," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 8697–8710.
- [6] H. Cai, J. Yang, W. Zhang, S. Han, and Y. Yu, "Path-level network transformation for efficient architecture search," in *Proceedings of the 35th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, J. Dy and A. Krause, Eds., vol. 80. PMLR, 10–15 Jul 2018, pp. 678–687. [Online]. Available: <https://proceedings.mlr.press/v80/cai18a.html>
- [7] G. Huang, Z. Liu, L. van der Maaten, and K. Q. Weinberger, "Densely connected convolutional networks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, July 2017.
- [8] Y. Liu, Y. Sun, B. Xue, M. Zhang, G. G. Yen, and K. C. Tan, "A survey on evolutionary neural architecture search," *IEEE transactions on neural networks and learning systems*, 2021.
- [9] A. Zela, A. Klein, S. Falkner, and F. Hutter, "Towards automated deep learning: Efficient joint neural architecture and hyperparameter search," *arXiv preprint arXiv:1807.06906*, 2018.
- [10] A. Klein, E. Christiansen, K. Murphy, and F. Hutter, "Towards reproducible neural architecture and hyperparameter search," 2018.
- [11] P. Chrabaszcz, I. Loshchilov, and F. Hutter, "A downsampled variant of imagenet as an alternative to the cifar datasets," *arXiv preprint arXiv:1707.08819*, 2017.
- [12] T. Domhan, J. T. Springenberg, and F. Hutter, "Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves," in *Twenty-fourth international joint conference on artificial intelligence*, 2015.
- [13] H. Cai, T. Chen, W. Zhang, Y. Yu, and J. Wang, "Efficient architecture search by network transformation," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, no. 1, 2018.
- [14] H. Liu, K. Simonyan, and Y. Yang, "Darts: Differentiable architecture search," *arXiv preprint arXiv:1806.09055*, 2018.
- [15] Y. Sun, G. G. Yen, and Z. Yi, "Igd indicator-based evolutionary algorithm for many-objective optimization problems," *IEEE Transactions on Evolutionary Computation*, vol. 23, no. 2, pp. 173–187, 2018.
- [16] L. Xie and A. Yuille, "Genetic cnn," in *Proceedings of the IEEE international conference on computer vision*, 2017, pp. 1379–1388.
- [17] S. Fujino, N. Mori, and K. Matsumoto, "Deep convolutional networks for human sketches by means of the evolutionary deep learning," in *2017 Joint 17th World Congress of International Fuzzy Systems Association and 9th International Conference on Soft Computing and Intelligent Systems (IFSA-SCIS)*. IEEE, 2017, pp. 1–5.
- [18] Y. Sun, B. Xue, M. Zhang, and G. G. Yen, "Evolving deep convolutional neural networks for image classification," *IEEE Transactions on Evolutionary Computation*, vol. 24, no. 2, pp. 394–407, 2019.
- [19] —, "Completely automated cnn architecture design based on blocks," *IEEE transactions on neural networks and learning systems*, vol. 31, no. 4, pp. 1242–1254, 2019.
- [20] Z. Chen, Y. Zhou, and Z. Huang, "Auto-creation of effective neural network architecture by evolutionary algorithm and resnet for image classification," in *2019 IEEE international conference on systems, man and cybernetics (SMC)*. IEEE, 2019, pp. 3895–3900.
- [21] D. Song, C. Xu, X. Jia, Y. Chen, C. Xu, and Y. Wang, "Efficient residual dense block search for image super-resolution," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 07, 2020, pp. 12 007–12 014.
- [22] E. Real, A. Aggarwal, Y. Huang, and Q. V. Le, "Regularized evolution for image classifier architecture search," in *Proceedings of the aaai*

- conference on artificial intelligence*, vol. 33, no. 01, 2019, pp. 4780–4789.
- [23] W. Irwin-Harris, Y. Sun, B. Xue, and M. Zhang, “A graph-based encoding for evolutionary convolutional neural network architecture design,” in *2019 IEEE Congress on Evolutionary Computation (CEC)*. IEEE, 2019, pp. 546–553.
 - [24] T. Elsken, J.-H. Metzen, and F. Hutter, “Simple and efficient architecture search for convolutional neural networks,” *arXiv preprint arXiv:1711.04528*, 2017.
 - [25] W. Zhang, L. Zhao, Q. Li, S. Zhao, Q. Dong, X. Jiang, T. Zhang, and T. Liu, “Identify hierarchical structures from task-based fmri data via hybrid spatiotemporal neural architecture search net,” in *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer, 2019, pp. 745–753.
 - [26] H. Zhu, Z. An, C. Yang, K. Xu, E. Zhao, and Y. Xu, “Eena: Efficient evolution of neural architecture,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops*, 2019, pp. 0–0.
 - [27] J. Lin, W.-M. Chen, Y. Lin, C. Gan, S. Han *et al.*, “Mcnnet: Tiny deep learning on iot devices,” *Advances in Neural Information Processing Systems*, vol. 33, pp. 11 711–11 722, 2020.
 - [28] I. Radosavovic, R. P. Kosaraju, R. Girshick, K. He, and P. Dollár, “Designing network design spaces,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 10 428–10 436.
 - [29] M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S. Ghemawat, G. Irving, M. Isard *et al.*, “{TensorFlow}: A system for {Large-Scale} machine learning,” in *12th USENIX symposium on operating systems design and implementation (OSDI 16)*, 2016, pp. 265–283.
 - [30] L. Lai, N. Suda, and V. C.-N. Chandra, “Efficient neural network kernels for arm cortex-m cpus. arxiv 2018,” *arXiv preprint arXiv:1801.06601*.
 - [31] M. Tan, B. Chen, R. Pang, V. Vasudevan, M. Sandler, A. Howard, and Q. V. Le, “Mnasnet: Platform-aware neural architecture search for mobile,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 2820–2828.